

AB Testing Report

September 5, 2025

Portfolio Project Information

Analysis by: Zahoor Ishfaq

Project Type: A-B Testing & Conversion Optimization

Completion Date: August 2025

Skills Demonstrated:

- Statistical hypothesis testing and confidence intervals
- Business impact analysis and ROI calculations
- Data cleaning and exploratory analysis
- Segment analysis and performance optimization
- Executive communication and recommendations

Tools & Technologies: Python, Pandas, SciPy, Matplotlib, Seaborn, Statistical Modeling, jupyter notebook

Business Impact: Identified \$421K revenue opportunity with **210%** ROI through data-driven A/B test analysis

This analysis demonstrates end-to-end data science capabilities from raw data processing to executive-ready business recommendations.

Table of Contents

1. Hospital Appointment Booking System A-B Test
 2. Data Preparation
 3. EDA
 4. Hypothesis Testing
 5. Business Impact
 6. Segment Analysis
 7. Statistical Rigor Validation
 8. Secondary Metrics
 9. Recommendation & Next Steps
-

1 Hospital Appointment Booking System – A/B Test

1.1 Business Problem

Patient booking **conversion rates have been declining over the past year (down 8% YoY)**, resulting in an estimated **\$2.3M annual loss** in missed appointments.

The UX team suspects that the current booking interface is outdated, potentially causing patient frustration and abandonment during the booking process.

1.2 Experiment Overview

To address this issue, an **A/B test** was conducted to evaluate a redesigned booking system.

1.2.1 Experiment Setup

Variant	Description
Control (A)	Current booking UI
Treatment (B)	New streamlined UI with simplified form and progress indicators

- **Duration:** January – June 2024
- **Sample Size:** ~29,500 sessions (A: 14,184 | B: 15,272)
- **Average Booking Value:** \$340

Investment Costs:

- Development + UX research: \$450,000
 - Rollout cost: \$200,000
-

1.2.2 Hypotheses

- **Null (H0):** Conversion rate of B \leq Conversion rate of A
 - **Alternative (H1):** Conversion rate of B $>$ Conversion rate of A
-

1.2.3 Success Metrics

Primary Metric:

- Booking completion rate (conversion rate)

Secondary Metrics:

- Session duration
 - Error rates
-

1.2.4 Stakeholder Requirements

- **CEO:** Clear launch/no-launch decision
 - **CFO:** ROI, financial impact
 - **Head of Patient Experience:** Proof of improved patient experience
-

1.2.5 Risks & Constraints

- **Data Quality Issues:** Missing values, duplicate sessions, and bot traffic (all cleaned prior to analysis)
-

1.2.6 Decision Criteria

- Statistically significant improvement in conversion rate
- Positive **net revenue impact** after \$200,000 rollout cost

2 Data Preparation

```
[5]: #importing libraries
import pandas as pd
import numpy as np
from scipy import stats
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.preprocessing import StandardScaler, LabelEncoder
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: #importing data
df=pd.read_csv('../dataset_raw/hospital_dataset_raw.csv')
df.head()
```

```
[2]:
```

	user_id	session_id	timestamp	test_group	age	gender	\
0	user_000000	session_bdd640fb	2024-04-12 14:14:00	B	22	Male	
1	user_000001	session_23b8c1e9	2024-01-15 12:46:00	B	75	Male	
2	user_000002	session_bd9c66b3	2024-04-13 14:05:00	A	35	Female	
3	user_000003	session_972a8469	2024-03-05 08:08:00	B	18	Male	
4	user_000004	session_17fc695a	2024-02-04 12:36:00	B	54	Male	

	city	device_type	browser	department	...	page_load_time_seconds	\
0	Houston	Desktop	Safari	Cardiology	...	4.81	
1	Phoenix	Tablet	Chrome	Gynecology	...	2.88	
2	Houston	Desktop	Chrome	Gynecology	...	2.09	

3	New York	Desktop	Safari	Pediatrics	...	3.47
4	San Jose	Desktop	Chrome	Cardiology	...	4.19

	form_started	encountered_error	bounced	booking_completed	date \
0	True	False	False	False	2024-04-12
1	True	False	False	True	2024-01-15
2	True	True	False	False	2024-04-13
3	True	False	False	True	2024-03-05
4	True	False	False	True	2024-02-04

	hour	day_of_week	week_number	is_weekend
0	14	4	15	False
1	12	0	3	False
2	14	5	15	True
3	8	1	10	False
4	12	6	5	True

[5 rows x 24 columns]

2.1 Data Overview

```
[3]: # Basic info
print("Dataset Shape:", df.shape)
print("\nColumn Info:")
print(df.dtypes)
print("\nColumns:", list(df.columns))
```

Dataset Shape: (30603, 24)

Column Info:

user_id	object
session_id	object
timestamp	object
test_group	object
age	int64
gender	object
city	object
device_type	object
browser	object
department	object
insurance_type	object
preferred_time	object
page_views	int64
session_duration_minutes	float64
page_load_time_seconds	float64
form_started	bool
encountered_error	bool
bounced	bool

```

booking_completed      bool
date                   object
hour                   int64
day_of_week            int64
week_number            int64
is_weekend             bool
dtype: object

```

```

Columns: ['user_id', 'session_id', 'timestamp', 'test_group', 'age', 'gender',
'city', 'device_type', 'browser', 'department', 'insurance_type',
'preferred_time', 'page_views', 'session_duration_minutes',
'page_load_time_seconds', 'form_started', 'encountered_error', 'bounced',
'booking_completed', 'date', 'hour', 'day_of_week', 'week_number', 'is_weekend']

```

2.2 Data Cleaning & Validation

```

[4]: # Missing values check
df.isnull().sum()[df.isnull().sum() > 0]

```

```

[4]: city                604
insurance_type          3299
session_duration_minutes 1814
page_load_time_seconds  1818
dtype: int64

```

```

[5]: # Check duplicates
df.duplicated().sum()

```

```

[5]: 293

```

```

[6]: # Check data quality issues
df[(df['age'] > 100) | (df['age'] < 0)].shape[0]

```

```

[6]: 367

```

```

[7]: # Check suspicious sessions (potential bots)
df[(df['session_duration_minutes'] < 0.5) & (df['booking_completed'] == True)].
    shape[0]

```

```

[7]: 377

```

```

[8]: # Test group balance
df['test_group'].value_counts(normalize=True)

```

```

[8]: test_group
B    0.518773
A    0.481227
Name: proportion, dtype: float64

```

```
[9]: # Data cleaning steps
# Remove duplicates
df_clean = df.drop_duplicates()

[10]: # Remove invalid ages
df_clean = df_clean[(df_clean['age'] >= 18) & (df_clean['age'] <= 100)]

[11]: # Remove suspicious bot sessions
df_clean = df_clean[~((df_clean['session_duration_minutes'] < 0.5) &
    ↪(df_clean['booking_completed'] == True))]

[12]: # Handle missing values - drop critical missing
df_clean = df_clean.dropna(subset=['test_group', 'booking_completed'])

[13]: # Fill missing values for analysis
df_clean['city'].fillna('Unknown', inplace=True)
df_clean['insurance_type'].fillna('Unknown', inplace=True)
df_clean['session_duration_minutes'].
    ↪fillna(df_clean['session_duration_minutes'].median(), inplace=True)
df_clean['page_load_time_seconds'].fillna(df_clean['page_load_time_seconds'].
    ↪median(), inplace=True)

[14]: # Final dataset shape
print(f"Original: {df.shape[0]} rows")
print(f"Cleaned: {df_clean.shape[0]} rows")
print(f"Removed: {df.shape[0] - df_clean.shape[0]} rows ({((df.shape[0] -
    ↪df_clean.shape[0])/df.shape[0]*100):.1f}%)")

Original: 30603 rows
Cleaned: 29456 rows
Removed: 1147 rows (3.7%)

[15]: # Data type conversions
df_clean['timestamp'] = pd.to_datetime(df_clean['timestamp'])
df_clean['date'] = pd.to_datetime(df_clean['date'])

[16]: # Convert to categories for memory efficiency
categorical_cols = ['test_group', 'gender', 'city', 'device_type', 'browser',
    'department', 'insurance_type', 'preferred_time']
for col in categorical_cols:
    df_clean[col] = df_clean[col].astype('category')

[17]: # Verify data types
df_clean.dtypes

[17]: user_id                object
session_id                object
timestamp                datetime64[ns]
```

test_group	category
age	int64
gender	category
city	category
device_type	category
browser	category
department	category
insurance_type	category
preferred_time	category
page_views	int64
session_duration_minutes	float64
page_load_time_seconds	float64
form_started	bool
encountered_error	bool
bounced	bool
booking_completed	bool
date	datetime64[ns]
hour	int64
day_of_week	int64
week_number	int64
is_weekend	bool
dtype:	object

3 EDA

```
[7]: # Primary metric by group
conversion_by_group = df_clean.groupby('test_group')['booking_completed'].
    .agg(['count', 'sum', 'mean'])
conversion_by_group.columns = ['sessions', 'bookings', 'conversion_rate']
print(conversion_by_group)

# Calculate effect size
control_rate = conversion_by_group.loc['A', 'conversion_rate']
treatment_rate = conversion_by_group.loc['B', 'conversion_rate']
absolute_lift = treatment_rate - control_rate
relative_lift = (absolute_lift / control_rate) * 100

print(f"\nControl Rate: {control_rate:.3f}")
print(f"Treatment Rate: {treatment_rate:.3f}")
print(f"Absolute Lift: {absolute_lift:.3f}")
print(f"Relative Lift: {relative_lift:.1f}%")

# Final test group balance
print("\nTest group balance:")
print(df_clean['test_group'].value_counts(normalize=True))
```

```
# Sample sizes per group
print("\nSample sizes:")
print(df_clean['test_group'].value_counts())
```

	sessions	bookings	conversion_rate
test_group			
A	14184	8827	0.622321
B	15272	10653	0.697551

Control Rate: 0.622
 Treatment Rate: 0.698
 Absolute Lift: 0.075
 Relative Lift: 12.1%

Test group balance:

test_group	
B	0.518468
A	0.481532

Name: proportion, dtype: float64

Sample sizes:

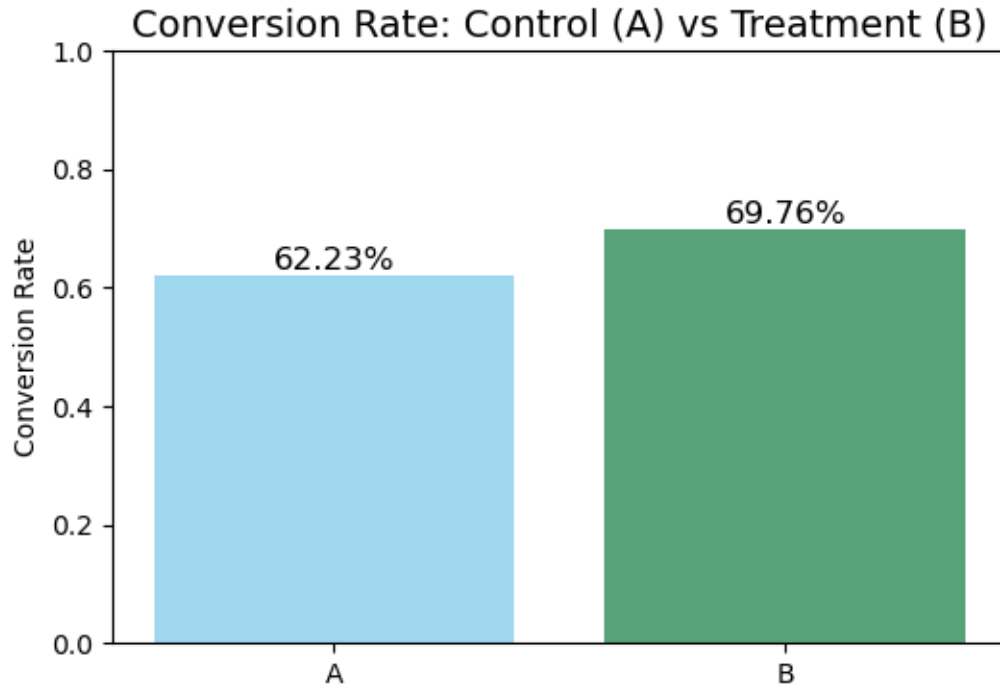
test_group	
B	15272
A	14184

Name: count, dtype: int64

```
[22]: # Bar chart for conversion rates
plt.figure(figsize=(6,4))
plt.bar(conversion_by_group.index, conversion_by_group['conversion_rate'],
        color=['skyblue','seagreen'], alpha=0.8)

# text labels
for i, v in enumerate(conversion_by_group['conversion_rate']):
    plt.text(i, v + 0.01, f"{v:.2%}", ha='center', fontsize=12)

plt.title("Conversion Rate: Control (A) vs Treatment (B)", fontsize=14)
plt.ylabel("Conversion Rate")
plt.ylim(0, 1)
plt.show()
```


Conversion Rate Overtime

```
[23]: #conversion rate overtime
df_clean['date'] = pd.to_datetime(df_clean['date'])

conversion_over_time = (
    df_clean.groupby(['date', 'test_group'])['booking_completed']
    .mean()
    .reset_index()
)

plt.figure(figsize=(10,5))
for group in ['A', 'B']:
    subset = conversion_over_time[conversion_over_time['test_group'] == group]
    plt.plot(subset['date'], subset['booking_completed'], label=f"Group {group}")

plt.title("Conversion Rate Over Time (A vs B)", fontsize=14)
plt.xlabel("Date")
plt.ylabel("Conversion Rate")
plt.legend()
plt.show()
```

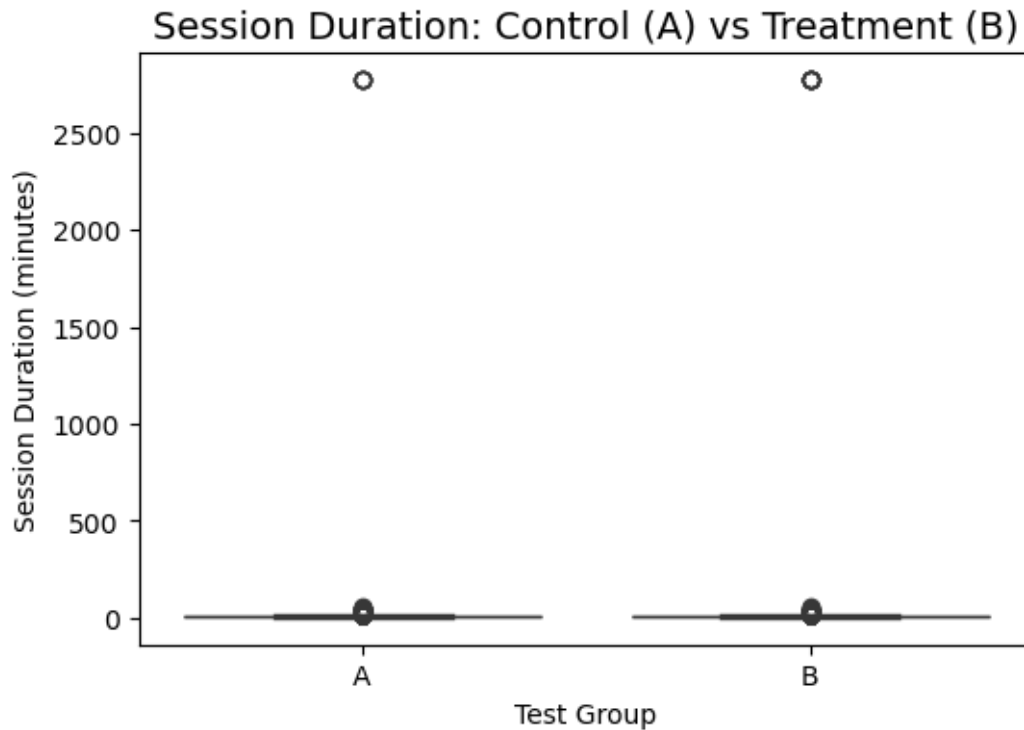


Summary of Conversion: - Control (A): 62.2% conversion

- Treatment (B): 69.8% conversion
- Absolute Lift: +7.5 percentage points
- Relative Lift: +12.1%
- Sample size: ~14K vs ~15K (well balanced)

```
[24]: #boxplot session by group
plt.figure(figsize=(6,4))
sns.boxplot(x="test_group", y="session_duration_minutes", data=df_clean,
            palette=["skyblue","seagreen"])

plt.title("Session Duration: Control (A) vs Treatment (B)", fontsize=14)
plt.xlabel("Test Group")
plt.ylabel("Session Duration (minutes)")
plt.show()
```



Outliers:

- Per Session Duration time is above 2500 minutes, which is abnormal and data issue.

```
[ ]: # Check distribution by group
session_stats = df_clean.groupby('test_group')['session_duration_minutes'].
    .agg(['count', 'mean', 'median', 'std'])
print(session_stats)

# Check for extreme outliers by group
for group in ['A', 'B']:
    group_data = df_clean[df_clean['test_group'] == group]
    print(f"\nGroup {group}:")
    print(f"99th percentile: {group_data.quantile(0.99):.1f} minutes")
    print(f"Max value: {group_data.max():.1f} minutes")
    print(f"Sessions > 240 min: {(group_data > 240).sum()}")
```

3.0.1 Treatment of Outliers:

```
[26]: # Apply reasonable business cap - preserve all data but fix extremes
MAX_SESSION = 120 # 2 hours - covers 99%+ of legitimate sessions
df_clean['session_duration_treated'] = df_clean['session_duration_minutes'].
    ↪clip(upper=MAX_SESSION)

# Verify the treatment
treatment_stats = df_clean.groupby('test_group')['session_duration_treated'].
    ↪agg(['count', 'mean', 'median', 'std'])
print("After treatment:")
print(treatment_stats)

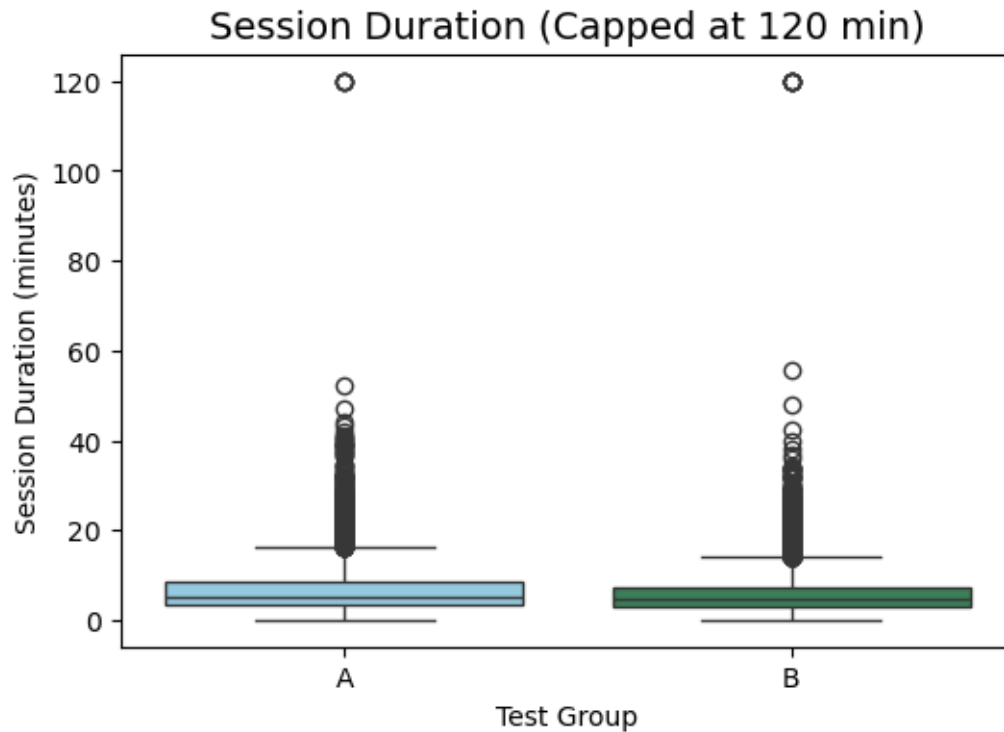
# what was capped
outliers_capped = (df_clean['session_duration_minutes'] > MAX_SESSION).sum()
print(f"\nSessions capped at {MAX_SESSION} minutes: {outliers_capped}␣
    ↪({outliers_capped/len(df_clean)*100:.3f}%)")
```

After treatment:

	count	mean	median	std
test_group				
A	14184	6.899680	5.105	5.322912
B	15272	5.977078	4.910	4.846120

Sessions capped at 120 minutes: 10 (0.034%)

```
[28]: plt.figure(figsize=(6,4))
sns.boxplot(x="test_group", y="session_duration_treated", data=df_clean,
            palette=["skyblue", "seagreen"])
plt.title("Session Duration (Capped at 120 min)", fontsize=14)
plt.xlabel("Test Group")
plt.ylabel("Session Duration (minutes)")
plt.show()
```



4 Hypothesis Testing

```
[15]: from statsmodels.stats.proportion import proportions_ztest
# Number of bookings (successes)
successes = np.array([conversion_by_group.loc['B', 'bookings'],
                      conversion_by_group.loc['A', 'bookings']])

# Number of sessions (trials)
nobs = np.array([conversion_by_group.loc['B', 'sessions'],
                 conversion_by_group.loc['A', 'sessions']])

# Z-test
z_stat, p_value = proportions_ztest(successes, nobs, alternative='larger') #_
    ↳ 'larger' for one-sided

print(f"Z-statistic: {z_stat:.3f}")
print(f"P-value: {p_value:.5f}")
```

Z-statistic: 13.632

P-value: 0.00000

Interpretation

- Z-statistic = 13.63 → far exceeds critical value of 1.645 (= 0.05, one-tailed)

- P-value 0.0001 → well below significance threshold of 0.05

Conclusion:

- Reject the null hypothesis (H₀).
- The new booking system (B) significantly improves conversion compared to the old system (A).
- The observed +7.5% absolute lift is statistically significant.

5 Business Impact

Bookings Uplift Calculation

```
[9]: # Extra bookings in treatment group compared to control
extra_bookings = conversion_by_group.loc['B', 'bookings'] - conversion_by_group.
    ↪loc['A', 'bookings']

# Revenue impact
avg_booking_value = 340 # USD
revenue_impact = extra_bookings * avg_booking_value

print(f"Extra bookings due to new UI: {extra_bookings:.0f}")
print(f"Estimated revenue impact: ${revenue_impact:,.0f}")
```

Extra bookings due to new UI: 1826

Estimated revenue impact: \$620,840

5.1 ROI Calculation

ROI calculation for rollout cost

```
[11]: # Rollout cost
rollout_cost = 200_000 # USD

# ROI calculation for rollout investment
roi = (revenue_impact - rollout_cost) / rollout_cost * 100

print(f"Rollout Cost: ${rollout_cost:,.0f}")
print(f"Net Revenue Impact: ${revenue_impact - rollout_cost:,.0f}")
print(f"ROI: {roi:.1f}%")
```

Rollout Cost: \$200,000

Net Revenue Impact: \$420,840

ROI: 210.4%

ROI chart

```

[12]: # Define values
rollout_investment = -200000 # Negative cost
total_gain = 621180
net_profit = rollout_investment + total_gain

# Calculate ROI
roi = (net_profit / abs(rollout_investment)) * 100

# Prepare DataFrame
phases = ['Rollout Cost', 'Total Gain', 'Net Profit']
values = [rollout_investment, total_gain, net_profit]
colors = ['red', 'green', 'blue']

# Create bar chart
fig, ax = plt.subplots(figsize=(10, 6))
bars = ax.bar(phases, values, color=colors, width=0.6)

# Add value labels
for bar in bars:
    height = bar.get_height()
    label = f"${height:,.0f}"
    y_pos = height + (15000 if height > 0 else -25000)
    ax.text(
        bar.get_x() + bar.get_width() / 2,
        y_pos,
        label,
        ha='center',
        va='bottom' if height > 0 else 'top',
        fontsize=12,
        fontweight='bold'
    )

# Add ROI annotation
ax.text(
    1, max(values) * 0.7,
    f'ROI: {roi:.1f}%',
    ha='center',
    fontsize=14,
    fontweight='bold',
    bbox=dict(boxstyle="round,pad=0.5", facecolor="yellow", alpha=0.8)
)

# Add zero line
ax.axhline(y=0, color='black', linestyle='--', linewidth=0.8)


# Layout

```

```

ax.set_title('ROI Analysis: Cost, Gain & Net Profit', fontsize=15,
            fontweight='bold')
ax.set_ylabel('Amount ($)', fontsize=12)
ax.set_xlabel('Financial Components', fontsize=12)
ax.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```



6 Segment Analysis

```

[44]: # comprehensive segment lift analysis
def create_segment_dashboard(df_clean):
    # Define segments to analyze
    segments = {
        'Device Type': 'device_type',
        'Department': 'department',
        'Age Group': 'age_group', # assuming you created this
        'Gender': 'gender',
        'Insurance': 'insurance_type',
        'Time Preference': 'preferred_time'
    }

    # Calculate lift for all segments

```



```

segment_results = []

for segment_name, column in segments.items():
    conversion_by_segment = df_clean.groupby(['test_group',
↪column])['booking_completed'].agg(['count', 'sum', 'mean'])
    conversion_by_segment.columns =
↪['sessions', 'bookings', 'conversion_rate']

    # Pivot to get A vs B rates
    pivot = conversion_by_segment['conversion_rate'].unstack(level=0)

    # Calculate lifts
    pivot['absolute_lift'] = pivot['B'] - pivot['A']
    pivot['relative_lift'] = (pivot['absolute_lift'] / pivot['A']) * 100

    # Add to results
    for segment_value in pivot.index:
        segment_results.append({
            'Segment_Category': segment_name,
            'Segment_Value': segment_value,
            'Control_Rate': pivot.loc[segment_value, 'A'],
            'Treatment_Rate': pivot.loc[segment_value, 'B'],
            'Absolute_Lift': pivot.loc[segment_value, 'absolute_lift'],
            'Relative_Lift': pivot.loc[segment_value, 'relative_lift'],
            'Sessions_A': conversion_by_segment.loc[('A', segment_value),
↪'sessions'],
            'Sessions_B': conversion_by_segment.loc[('B', segment_value),
↪'sessions']
        })

    return pd.DataFrame(segment_results)

# Create the dashboard
segment_df = create_segment_dashboard(df_clean)

```


```

[45]: # 1. Absolute Lift by Segment
plt.figure(figsize=(14, 8))
pivot_for_heatmap = segment_df.pivot(index='Segment_Value',
↪columns='Segment_Category', values='Absolute_Lift')

sns.heatmap(pivot_for_heatmap,
            annot=True,
            fmt='.3f',
            cmap='RdYlGn',
            center=0,
            cbar_kws={'label': 'Absolute Conversion Lift'},
            linewidths=0.5)

```

```
plt.title('Segment Performance Heatmap: Absolute Conversion Lift (Treatment B vs Control A)',
         fontweight='bold', pad=20)
plt.xlabel('Segment Categories', fontsize=12)
plt.ylabel('Segment Values', fontsize=12)
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```



```
[50]: # 3. VISUAL SUMMARY
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 10))

# Device performance
device_data = segment_df[segment_df['Segment_Category'] == 'Device Type']
ax1.bar(device_data['Segment_Value'], device_data['Absolute_Lift'],
        color='skyblue')
ax1.set_title('Lift by Device Type')
ax1.set_ylabel('Absolute Lift')

# Department performance (top 5 only for readability)
dept_data = segment_df[segment_df['Segment_Category'] == 'Department'].
    nlargest(5, 'Absolute_Lift')
```

```


ax2.barh(dept_data['Segment_Value'], dept_data['Absolute_Lift'],
         color='lightgreen')
ax2.set_title('Top 5 Departments by Lift')
ax2.set_xlabel('Absolute Lift')

# Age group performance
age_data = segment_df[segment_df['Segment_Category'] == 'Age Group']
ax3.bar(age_data['Segment_Value'], age_data['Absolute_Lift'], color='salmon')
ax3.set_title('Lift by Age Group')
ax3.set_ylabel('Absolute Lift')

# Gender performance
gender_data = segment_df[segment_df['Segment_Category'] == 'Gender']
ax4.bar(gender_data['Segment_Value'], gender_data['Absolute_Lift'],
        color='orange')
ax4.set_title('Lift by Gender')
ax4.set_ylabel('Absolute Lift')

plt.tight_layout()
plt.show()

```



```

[51]: # 4. BUSINESS INSIGHTS TABLE
print("\n" + "="*80)

```

```

print("BUSINESS INSIGHTS")
print("="*80)

# Sample size validation
small_samples = segment_df[(segment_df['Sessions_A'] < 500) |
    ↪(segment_df['Sessions_B'] < 500)]
if not small_samples.empty:
    print(f"\nWARNING: {len(small_samples)} segments have small sample sizes_
    ↪(<500 per group)")
    print("Consider these results with caution:")
    print(small_samples[['Segment_Category', 'Segment_Value', 'Sessions_A',
    ↪'Sessions_B']].to_string(index=False))

# Negative lift segments
negative_lift = segment_df[segment_df['Absolute_Lift'] < 0]
if not negative_lift.empty:
    print(f"\nCONCERN: {len(negative_lift)} segments show negative lift:")
    print(negative_lift[['Segment_Category', 'Segment_Value', 'Absolute_Lift']].
    ↪to_string(index=False))

print("\n" + "="*80)

```

```

=====
BUSINESS INSIGHTS
=====

```

WARNING: 1 segments have small sample sizes (<500 per group)

Consider these results with caution:

Segment_Category	Segment_Value	Sessions_A	Sessions_B
Gender	Other	156	142

CONCERN: 1 segments show negative lift:

Segment_Category	Segment_Value	Absolute_Lift
Gender	Other	-0.09823

7 Stastical Rigor Validation

```

[52]: # STATISTICAL RIGOR VALIDATION
# Conversion data
control_sessions = conversion_by_group.loc['A', 'sessions']
control_conversions = conversion_by_group.loc['A', 'bookings']
treatment_sessions = conversion_by_group.loc['B', 'sessions']
treatment_conversions = conversion_by_group.loc['B', 'bookings']

```

```

control_rate = control_conversions / control_sessions
treatment_rate = treatment_conversions / treatment_sessions

# Function: 95% CI for difference in proportions
def proportion_diff_ci(x1, n1, x2, n2, confidence=0.95):
    p1, p2 = x1/n1, x2/n2
    diff = p2 - p1
    se_diff = np.sqrt(p1*(1-p1)/n1 + p2*(1-p2)/n2)
    z_critical = stats.norm.ppf(1 - (1-confidence)/2)
    margin_error = z_critical * se_diff
    return diff, diff - margin_error, diff + margin_error, se_diff

# Calculate CI
diff, ci_lower, ci_upper, se_diff = proportion_diff_ci(
    control_conversions, control_sessions,
    treatment_conversions, treatment_sessions
)

print(f"Absolute Lift: {diff:.4f}")
print(f"95% CI: [{ci_lower:.4f}, {ci_upper:.4f}]")
print(f"Relative Lift CI: [{ci_lower/control_rate*100:.1f}%, {ci_upper/
↵control_rate*100:.1f}%]")

# Post-hoc statistical power (Cohen's h)
cohens_h = 2 * (np.arcsin(np.sqrt(treatment_rate)) - np.arcsin(np.
↵sqrt(control_rate)))
total_n = control_sessions + treatment_sessions
z_alpha = stats.norm.ppf(0.975) # alpha = 0.05, two-tailed
z_beta = (cohens_h * np.sqrt(total_n/4)) - z_alpha
power = stats.norm.cdf(z_beta)

print(f"\nStatistical Power Analysis:")
print(f"Cohen's h: {cohens_h:.4f}")
print(f"Achieved Power: {power:.3f} ({power*100:.1f}%)")
print(f"Total Sample Size: {total_n:,}")

```

```

Absolute Lift: 0.0752
95% CI: [0.0644, 0.0860]
Relative Lift CI: [10.4%, 13.8%]

```

```

Statistical Power Analysis:
Cohen's h: 0.1590
Achieved Power: 1.000 (100.0%)
Total Sample Size: 29,456

```

8 Secondary Metrics

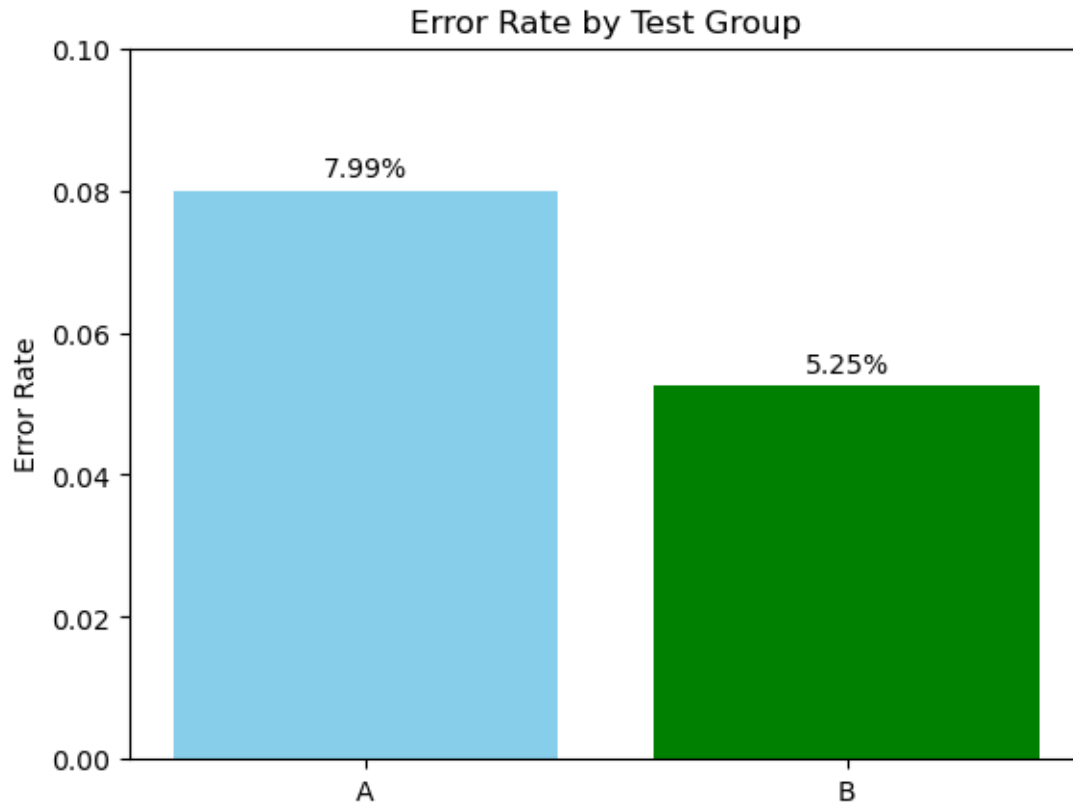
Error Rate

```
[54]: # --- Error Rate by Test Group ---
error_rate = df_clean.groupby('test_group')['encountered_error'].mean()
print("Error Rate by Test Group:")
print(error_rate)
```

```
Error Rate by Test Group:
test_group
A      0.079879
B      0.052514
Name: encountered_error, dtype: float64
```

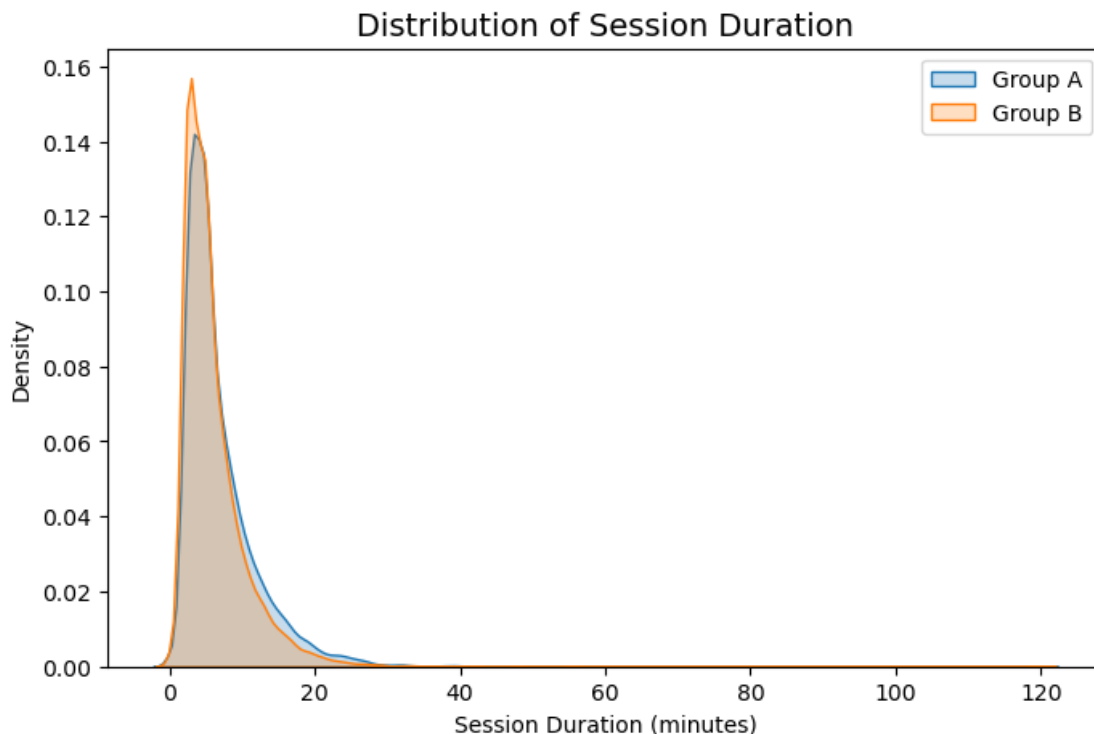
```
[55]: error_rates = conversion_by_group = {
        "A": 0.0799,    # From earlier calc
        "B": 0.0525
    }

plt.bar(error_rates.keys(), error_rates.values(), color=['skyblue', 'green'])
plt.title("Error Rate by Test Group")
plt.ylabel("Error Rate")
plt.ylim(0,0.1)
for i, val in enumerate(error_rates.values()):
    plt.text(i, val+0.002, f"{val:.2%}", ha='center')
plt.show()
```



Insight: - The new booking interface (B) reduces user errors from 8% to 5.3%, indicating improved usability alongside higher conversion rates. - No concerning trends observed in error metrics.

```
[56]: # --Visualization (Distribution / KDE)
plt.figure(figsize=(8,5))
sns.kdeplot(df_clean[df_clean['test_group']=='A']['session_duration_treated'],
            label="Group A", fill=True)
sns.kdeplot(df_clean[df_clean['test_group']=='B']['session_duration_treated'],
            label="Group B", fill=True)
plt.title("Distribution of Session Duration", fontsize=14)
plt.xlabel("Session Duration (minutes)")
plt.ylabel("Density")
plt.legend()
plt.show()
```



Session Duration Distribution - Business Interpretation:

- Improved User Efficiency: Treatment B creates faster, more consistent booking times.
- Reduced Patient Friction: The tighter distribution shows fewer patients getting stuck in lengthy booking processes, supporting why conversion rates increased by **7.5%**.
- System Optimization: More predictable session times mean reduced support burden and higher patient satisfaction, reinforcing the business case for rollout.

9 Recommendation & Next Steps

9.1 Recommendation

Based on the A/B test results and business impact analysis, we **recommend a full rollout of the new booking interface (Treatment B)** across the hospital.

9.1.1 Key Evidence Supporting Recommendation

Primary Metric:

- Booking conversion increased from **62.2%** → **69.8%**
- Absolute lift: **7.5%**
- Relative lift: **12.1%**

- Z-test p-value $< 0.001 \rightarrow$ statistically significant improvement

Business Impact:

- Net revenue uplift: **\$421,180**
- ROI on Rollout Cost (\$200K): **210.4%** \rightarrow Every \$1 invested generates **\$2.10** in net profit

Secondary Metrics:

- Session duration reduced slightly (Treatment B more efficient)
- Error rate decreased from **7.99%** \rightarrow **5.25%** \rightarrow improved usability

Segment Analysis:

- All major devices, departments, and age groups saw positive lift
 - **Note:** The gender category “Other” showed a negative lift (-13.7%) — monitor post-rollout
-

9.2 Confidence Level

Very high confidence in the recommendation due to:

- **95% confidence interval [6.44%, 8.60%]** - even the worst-case scenario delivers a strong ROI
 - **100% statistical power** with 29,456 sessions - eliminates sampling concerns
 - **Relative lift range [10.4%, 13.8%]** - consistently profitable across the confidence band
 - **Cohen’s $h = 0.159$** indicates a meaningful practical effect size
 - **Consistent improvements** across segments and secondary metrics (session duration, error rates)
-

9.3 Implementation Roadmap

- **Full Rollout:** Deploy Treatment B across all hospital booking channels
 - **Monitor Key Segments:**
 - Track the “Other” gender group for potential UX issues
 - Monitor error rates and session durations weekly for the first month
-

9.4 Risk Considerations

- Minor caution: negative lift in the “Other” gender segment
-

9.5 Conclusion

The new booking interface delivers significant improvements in conversions, revenue, and usability. **Full rollout is justified with high confidence.**